

This YouTube video by Krishak focuses on various text summarization techniques using Langchain and OpenAI's LLMs. Krishak emphasizes that while many believe such tasks require paid APIs, they can also be accomplished with open-source alternatives, though he uses OpenAI for greater accuracy in this demonstration. He promises a future video showcasing an end-to-end generative AI project applying these techniques.

The video details four to five text summarization methods. The first, "basic prompt summarization," uses a simple prompt directly with the LLM (ChatGPT 3.5 Turbo in this case). The process involves setting up system, human, and AI messages within the Langchain framework. The speaker demonstrates summarizing a speech by Narendra Modi, highlighting the use of ``system message``, ``human message``, and ``AI message`` to structure the interaction with the LLM. The code snippet demonstrates how to initialize the OpenAI API key, construct the prompt, send it to the LLM, and retrieve the summarized output. Token limits within LLMs are discussed, emphasizing the importance of considering token counts when working with large text inputs.

The second technique introduces prompt templates for more customized summarization. This involves creating a custom prompt using ``Langchain.prompt_template`` and ``llm_chain`` to handle multiple prompts sequentially. The example shows creating a template to summarize a speech and then translate the summary into Hindi. The code demonstrates how to create the template, format it with the input speech and desired language, and use the ``llm_chain`` to execute the summarization.

The third method, the "stuff documentation chain," involves feeding the entire document text to the LLM at once. This is efficient for smaller documents, but limitations arise when dealing with texts exceeding the LLM's token limit. The example uses a two-page PDF of a speech by A.P.J. Abdul Kalam. The code demonstrates reading the PDF, converting its content into a Langchain ``Document`` object, and using ``load_summarized_chain`` with ``chain_type="stuff"`` to perform the

summarization.

For larger documents, the video advocates for the "map reduce" technique. This method breaks down the document into smaller chunks, summarizes each chunk individually, and then combines the individual summaries into a final summary. The code demonstrates using ``recursive_character_text_splitter`` to divide the text into chunks, processing each through the LLM, and then combining the results. The concept of using custom prompts for both the individual chunk summaries and the final combined summary is also shown, enhancing control and customization.

Finally, the video briefly introduces the "refine" technique, similar to map reduce but iteratively refining the summary by incorporating progressively larger portions of the document. The presenter suggests this as another option for handling large documents and encourages viewers to experiment with custom prompts within this method. The video concludes with an assignment for viewers to implement refine with custom prompts and a preview of an upcoming project that will build a full application incorporating these text summarization techniques.